

NATIONAL UNIVERSITY OF SINGAPORE

SCHOOL OF COMPUTING

EXAMINATION FOR
Semester 1 AY2010/11

CS2100 – COMPUTER ORGANISATION

Nov 2010

Time allowed: 2 hours

Your Matriculation Number:

INSTRUCTIONS TO CANDIDATES

1. This examination paper consists of **SIX (6)** questions and comprises **TWENTY-ONE (21)** printed pages including this page.
2. This is an **OPEN BOOK** examination. You may use any approved calculators but not any PDA or laptop, especially those capable of external connectivity or communication.
3. Answer all questions. Note that the full mark for each question is different.
4. Write your answers on *this* **QUESTION AND ANSWER SCRIPT**. Answer only in the space given. Any writing outside this space will not be considered. No other submission is allowed.
5. Fill in your Matriculation Number with a pen, clearly on every page of this **QUESTION AND ANSWER SCRIPT**.
6. You may use pencil to write your answers.
7. At the end of the examination, please check to ensure that your script has all the pages properly stapled together.
8. Note that when a number is written as “0xNNNN” it means that “NNNN” is in base 16.

Total Score

/100

QUESTION 1 (15 marks)

- (1a) Encode the MIPS instruction “**lui \$5, 0xFEED**”, leaving your answer as an 8-hexadecimal digit number. (2 marks)

ANSWER:

0x3c05feed

- (1b) Encode the MIPS instruction “**beq \$t1, \$s0, Label**”, leaving your answer as an 8-hexadecimal digit number. Assume that the PC of this instruction is 0x1000 and Label is located at address 0x880. (2 marks)

ANSWER:

$$\text{Target_pc} = (\text{PC} + 4) + (4 * \text{Immed})$$

$$\text{Immed} = (0x880 - (0x1000 + 4)) / 4 = -0x1E1 = 0xFE1F \text{ (in 16 bits)}$$

0x1130fe1f

(1c) What is the MIPS instruction that is encoded by the hexadecimal number **0x00016C42**? (2 marks)

ANSWER:

```
srl $13, $1, 17
```

(1d) Convert following C code snippet into MIPS assembly.

```
sum = 0;
for (i=0; i<N; i++) {
    if (A[i] == i) continue;
    sum += A[i];
}
```

Note that: sum is in \$s0, address of the word (4-byte) array A is in \$a0, i is in \$t0, and N is in \$s1. (6 marks)

ANSWER:

```

        li      $s0, 0
        li      $t0, 0
Loop:   slt      $t1, $t0, $s1
        beq     $t1, $0, Exit
        sll     $t2, $t0, 2
        add     $t3, $t2, $a0
        lw      $t4, 0($t3)
        beq     $t4, $t0, Cont
        add     $s0, $s0, $t4
Cont:   addi     $t0, $t0, 1
        j       Loop
Exit:
```

(1e) Consider the following C code:

```
if (x > y) goto FARFARAWAY;
```

Suppose x is in $\$t0$, and y is in $\$t1$. The PC of this instruction is $0x100$ and `FARFARAWAY` is located at address $0xEEEE1800$. Give the MIPS instructions that will implement this `if`-statement. (3 marks)

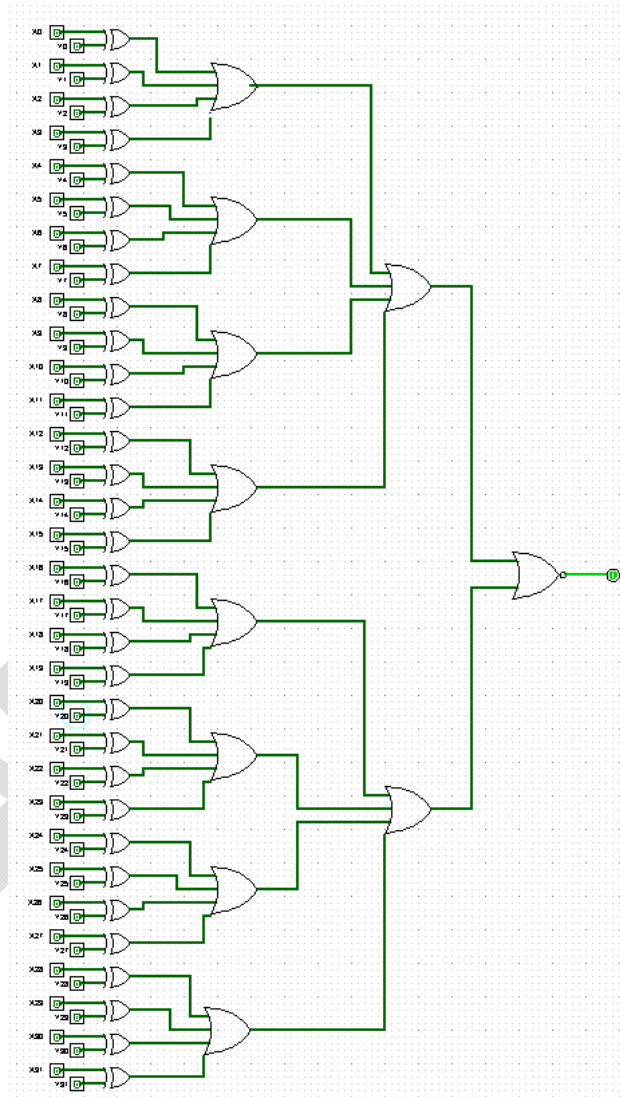
ANSWER:

```
        slt    $t2, $t1, $t0
        bne    $t2, $0, dojump
dojump:  j      FARFARAWAY
nojump:  ...
```

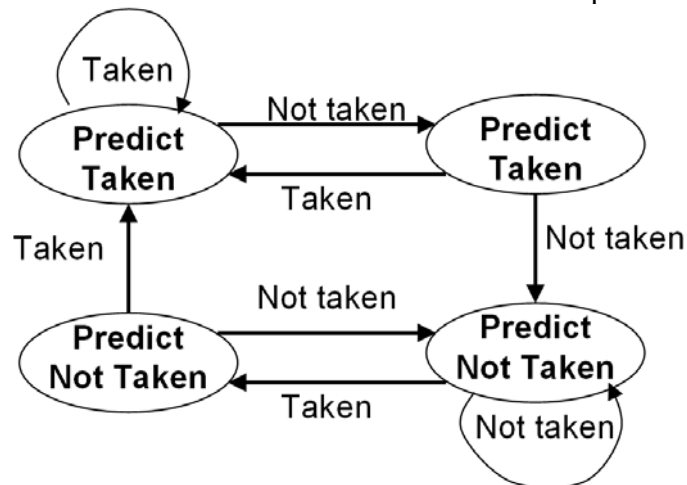
QUESTION 2 (15 marks)

- (2a) The MIPS beq/bne requires the comparison of two 32-bit quantities. If it is performed at the ALU, subtraction would be used. However, when we move it to the ID stage, we need a specialized comparator unit that compares the two quantities and outputs a single bit result indicating whether the two quantities are equal or not. There is no need for subtraction. Using basic gates, construct a circuit that will perform such a comparison. Assume that any gate has a maximum fan-in of 4 inputs. (5 marks)

ANSWER:

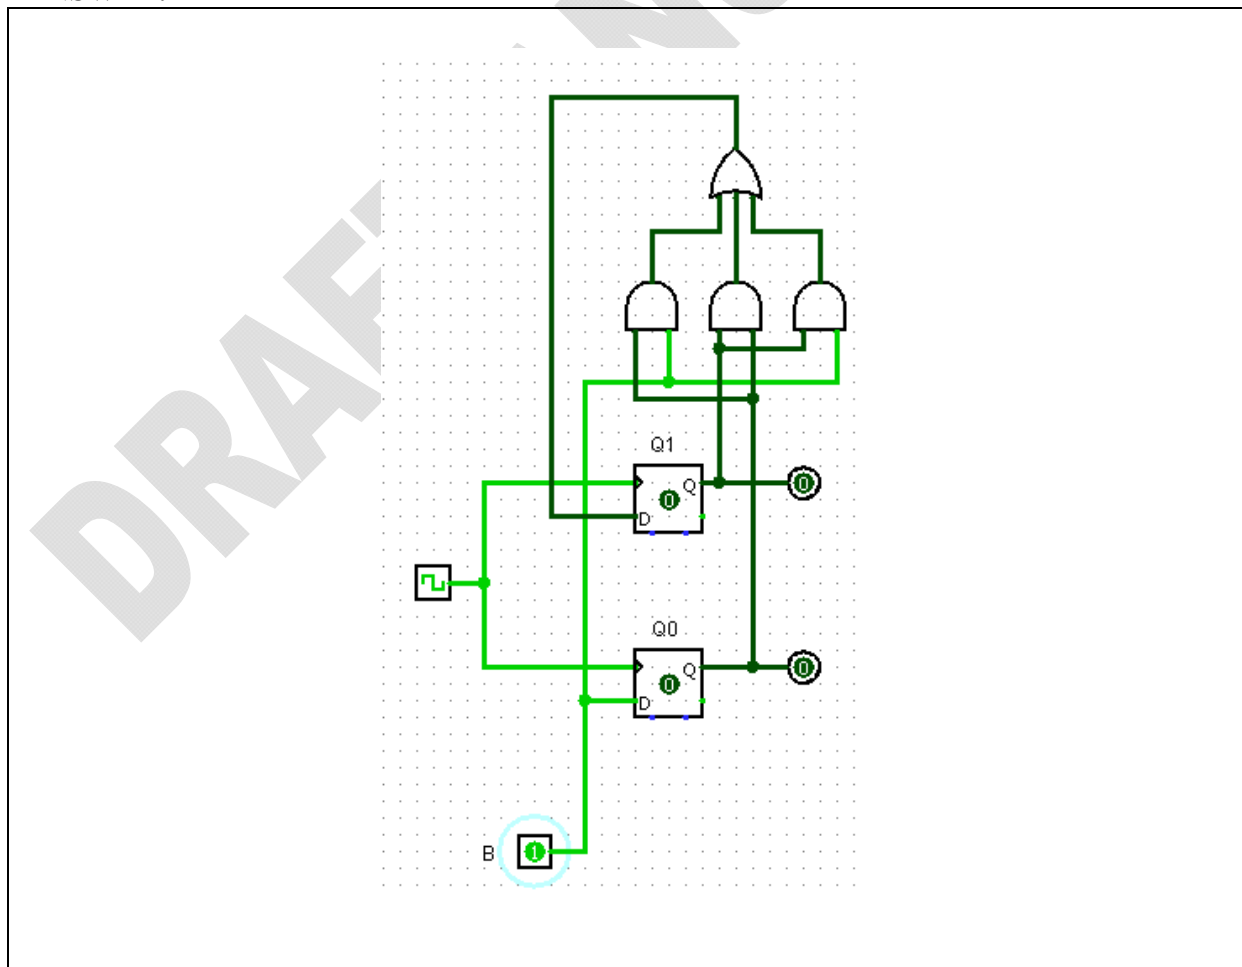


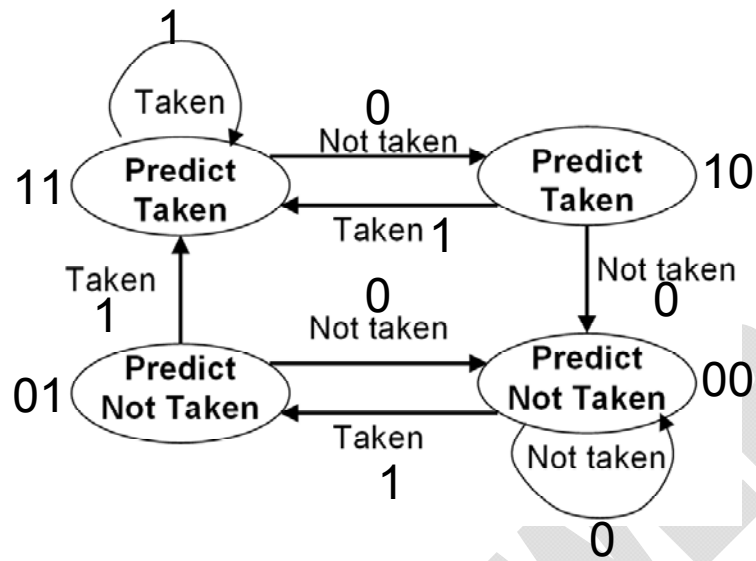
(2b) The following is the finite state machine for the 2-bit branch predictor.



Implement this finite state machine (including the prediction output) using two D-flip-flops, and the standard gates. The “taken/not-taken” (“0” = “not taken”, “1” = “taken”) is a single bit of input. The prediction, i.e., “Predict Taken/Not Taken” is a 1-bit output. Show clearly, how you encode the 4 states and your intermediate workings, including the truth table from which you derived the implementation. You may assume that the flip-flops have a fixed clock input. (10 marks)

ANSWER:





Q1	Q0	B	D1	D0
0	0	0	0	0
0	0	1	0	1
0	1	0	0	0
0	1	1	1	1
1	0	0	0	0
1	0	1	1	1
1	1	0	1	0
1	1	1	1	1

$$D1 = Q0 \cdot B + Q1 \cdot Q0 + Q1 \cdot B$$

$$D0 = B$$

QUESTION 3 (20 marks)

(3a) Given a NULL-terminated C string **str** (i.e., an array of characters terminated by an element that holds the value 0) that holds a string of at most 8 hexadecimal characters, write a MIPS assembly function “**int convert(str)**” that converts the hexadecimal string into an integer as the function’s return result. The address of **str** is passed as the first argument to **convert**. Note: the ASCII code for ‘0’ to ‘9’ is 48₁₀ to 57₁₀, ‘A’ to ‘F’ is 65₁₀ to 70₁₀, and ‘a’ to ‘f’ is 97₁₀ to 102₁₀, respectively. If any of the characters is not in these ranges, an error status is returned by returning a ‘1’ in \$v1. In this case, the content of \$v0 does not matter. If there is no error, then \$v0 should contain the converted value, and \$v1 should be ‘0’. You may use the various branching pseudo-ops for your convenience. (10 marks)

ANSWER:

```

convert:
    li    $v0, 0           // The return result
    li    $t0, 0           // The index
    li    $t4, 4           // The shift amount
    li    $t1, 8           // Limit
Loop:
    sll    $v0, $v0, $t4    // First iteration, $v0 = 0
    lw     $t2, 0($a0)      // str[i]
    li     $t3, 57          // 48 <= A[i] <= 57
    bgt    $t2, $t3, C65
    li     $t3, 48
    blt    $t2, $t3, Err
    addi   $t2, $t2, -48
    add    $v0, $v0, $t2
    j      Next
C65:
    li     $t3, 70          // 65 <= A[i] <= 70
    bgt    $t2, $t3, C97
    li     $t3, 65
    blt    $t2, $t3, Err
    addi   $t2, $t2, -55
    add    $v0, $v0, $t2
    j      Next
C97:
    li     $t3, 102         // 97 <= A[i] <= 102
    bgt    $t2, $t3, Err
    li     $t3, 97
    blt    $t2, $t3, Err
    addi   $t2, $t2, -87
    add    $v0, $v0, $t2
Next:
    addi   $t0, $t0, 1
    blt    $t0, $t1, Loop
    li     $v1, 0
Return:

```

jr	\$ra
li	\$v1, 1
j	Return

DRAFT ANSWERS

- (3b) The n factorial is defined as $\text{fac}(n) = n * \text{fac}(n-1)$ with $\text{fac}(1) = 1$, and $n > 0$. Write a MIPS assembly function that computes the n factorial where n is the input argument. (10 marks)

ANSWER:

```
.text
.globl main
# Prompt for a non-negative integer
# and invoke the factorial function.

main:
    addi $sp, $sp, -4    # Make space on stack.
    sw   $ra, 0($sp)    # Save return address.
    la   $a0, prompt
    li   $v0, 4
    syscall                # Display prompt.
    li   $v0, 5
    syscall                # Get integer response.
    move $a0, $v0        # Call factorial function.
    jal  factorial
    move $a0, $v0
    li   $v0, 1
    syscall                # Print integer result.
    la   $a0, endl
    li   $v0, 4
    syscall                # Print endl.
    li   $v0, 0           # Return zero.
    lw   $ra, 0($sp)      # Restore return address.
    addi $sp, $sp, 4      # Restore stack pointer.
    jr   $ra

factorial:
    addi $sp, $sp, -8    # Make space on stack.
    sw   $ra, 0($sp)    # Save return address.
    li   $v0, 1          # 0! = 1
    beqz $a0, zero      # Special case for 0!
    sw   $a0, 4($sp)     # Save our argument.
    addi $a0, $a0, -1    # Calculate (n-1)!
    jal  factorial       # Result in v0.
    lw   $a0, 4($sp)     # Restore our argument.
    mul  $v0, $a0, $v0    # n! = n * (n-1)!

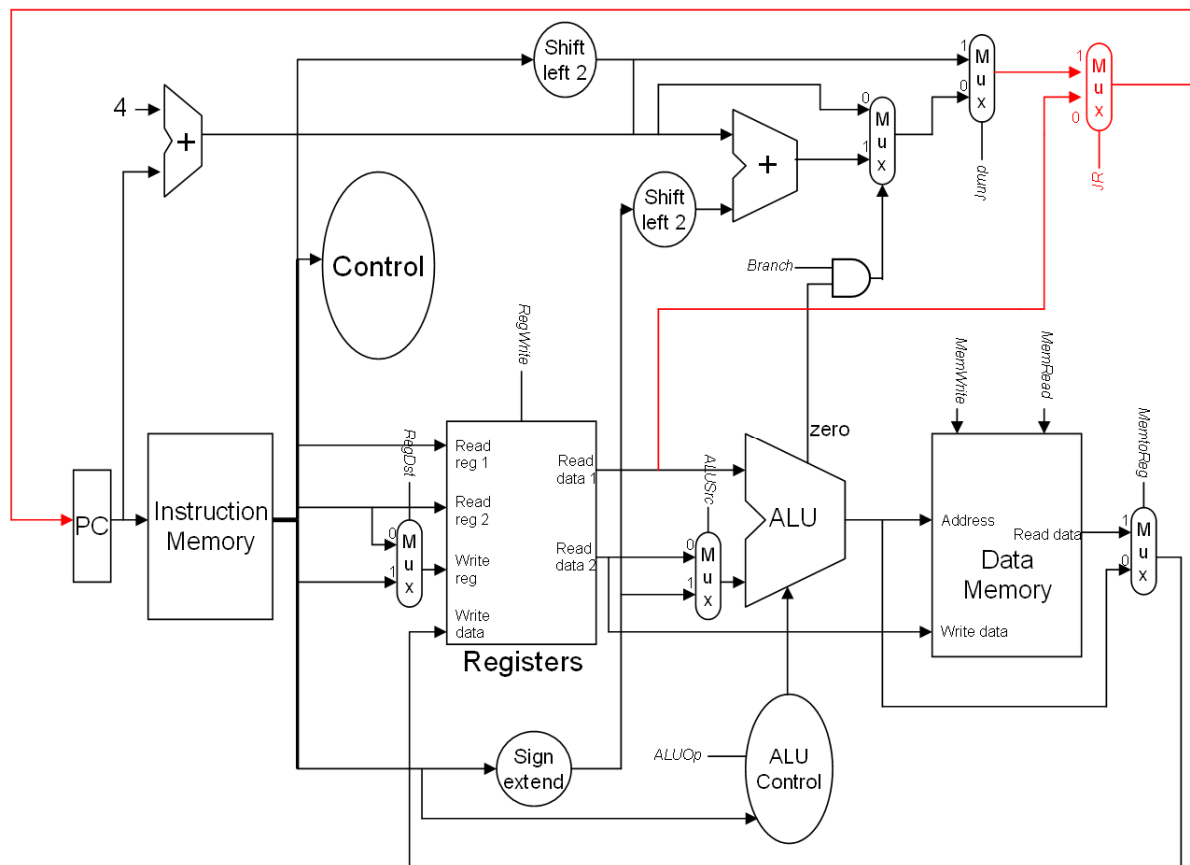
zero:
    lw   $ra, 0($sp)     # Restore return address.
    addi $sp, $sp, 8     # Restore stack pointer.
    jr   $ra            # Return.

.data
prompt: .asciiz "Enter a non-negative integer: "
endl:   .asciiz "\n"
```

QUESTION 4 (20 marks)

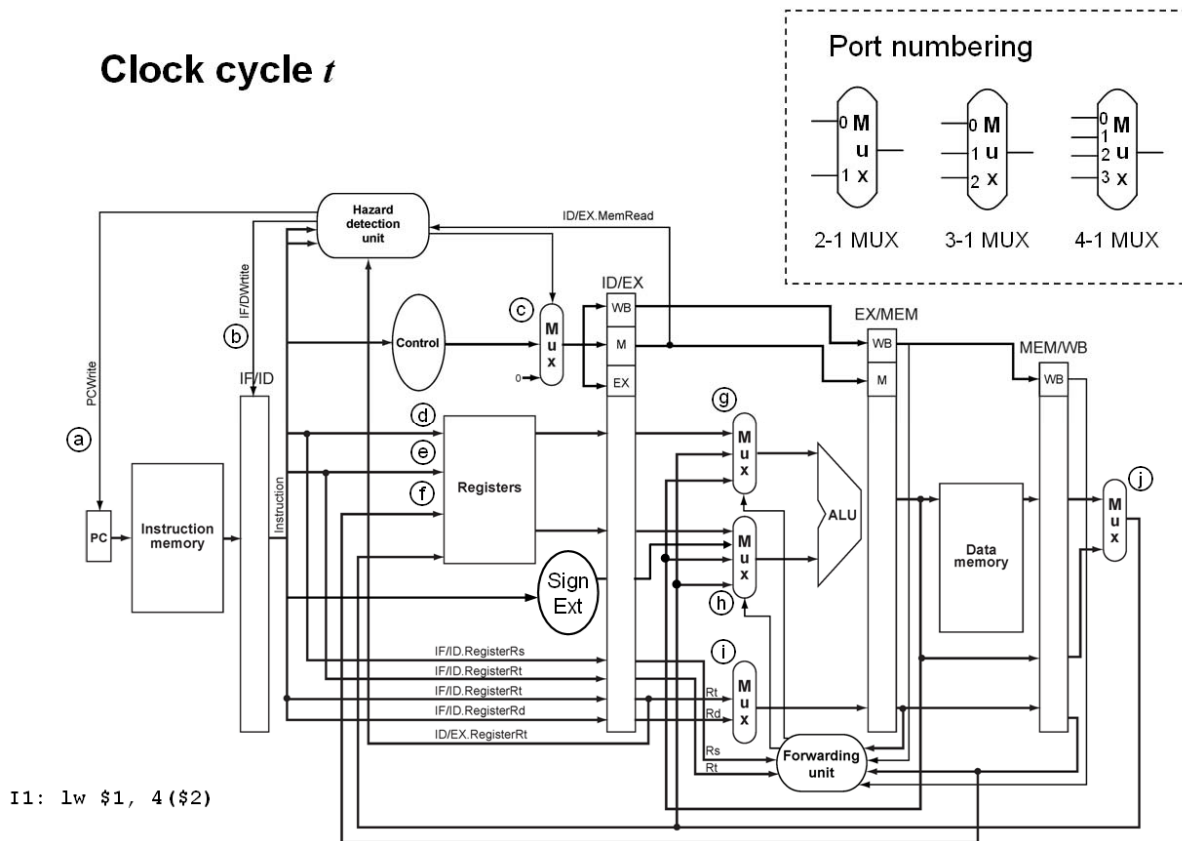
- (4a) The diagram below is a simplified version of the single cycle datapath. By drawing additional components and signal lines, including possibly new control signals (which you should define), show how support for the **jr** (“jump register”) instruction can be added. (10 marks)

ANSWER:



- (4b) Consider the following pipelined data path at clock cycle t . (The multiplexor's inputs are labeled as per the diagrams in the box with the dotted border.) The instruction **I1**, namely **lw \$1, 4(\$2)** is currently at the instruction fetch stage.

Clock cycle t



After **I1**, the following instructions enter the pipeline in order:

I2: add \$1, \$1, \$1

I3: sub \$2, \$1, \$5

I2 starts execution in the IF stage in cycle $t+1$, **I3** starts at cycle $t+2$, and so on. For the cycles, what are the values of the respective signals marked by the circled letters? Assume that the instructions executing/executed prior to cycle t has no effect over the current execution (other than the default $PC = PC+4$). (10 marks)

ANSWER:

Cycle	Signal	Value	Reason
$t+1$	(a)	1	Enable PC write
	(b)	1	Enable IF write
	(c)	0	Pass on control signal
	(d)	2	Rs of LW
	(e)	1	Rt of LW
	(f)	?	Unable to determine – depends on the previous instruction
$t+2$	(a)	0	Hazard detected. Introduce a stall.
	(b)	0	Hazard detected. Introduce a stall.
	(c)	1	No-op injected.
	(d)	1	Rs of ADD
	(e)	1	Rt of ADD
	(f)	?	Unable to determine – depends on the previous instruction
	(g)	0	Rs of LW
	(h)	1	Immediate of LW
	(i)	0	LW writes to Rt
$t+3$	(a)	1	Resume instruction fetch
	(b)	1	Resume instruction decode
	(c)	0	Control for ADD
	(d)	1	Rs of ADD
	(e)	1	Rt of ADD
	(f)	?	Unable to determine – depends on the previous instruction
	(g)	0	No-op
	(h)	0	No-op
	(i)	0	No-op
$t+4$	(a)	1	Enable PC write
	(b)	1	Enable IF write
	(c)	0	Pass on control signal of SUB
	(d)	1	Rs of SUB

	Ⓔ	5	Rt of SUB
	Ⓕ	1	Rt of LW
	Ⓖ	1	Forward from WB stage
	Ⓗ	3	Forward from WB stage
	Ⓙ	1	ADD writes to Rd
	⓫	0	LW data is from memory.
$t+5$	Ⓕ	0	NOP
	Ⓖ	2	Forward Rs value from MEM stage
	Ⓗ	0	Rt of SUB
	Ⓙ	1	Rd of SUB
	⓫	0	NOP
$t+6$	Ⓕ	1	Rd of ADD
	⓫	1	Pass on ALU result

QUESTION 5 (20 marks)

A machine with byte addresses and a word size of 32 bits and address width of 32 bits has a small 1024-block *direct-mapped* write-back cache with each block consisting of 2 words.

- (5a) Propose a C/Java data structure that can be used to hold all the content of the cache. You may assume that the data type “long” holds a 4-byte quantity. (4 marks)

ANSWER:

```
struct cache_block {
    long tag;          // Tag
    int valid_p;       // The valid bit
    int dirty_bit;     // For write-back
    long word[2];
} cache[1024];
```


- (5b)** Using C or Java, write a function that will take a 32 bit address as an argument (typed as “long”) and return a 32 bit data (again, a “long”) if it is found in the cache. It sets a global (integer) flag to indicate a hit (1) or a miss (0). You do not need to concern yourself with further miss processing in the lower cache hierarchy. (8 marks)

ANSWER:

```
int hit_flag = 0;

long read_cache(long addr) {
    int blkno, wordno;
    long tag;

    wordno = (addr >> 2) & 1;
    blkno = (addr >> 3) & 0x3FF;
    tag = addr >> 13;

    if (cache[blkno].tag == tag) { // hit
        hit_flag = 1;
        return(cache[blkno].word[wordno]);
    }
    else { // miss
        hit_flag = 0;
        return 0;
    }
}
```

- (5c) Using C or Java, write a code fragment will cause a lot of *compulsory* misses in this cache. (8 marks)

ANSWER:

The following loop will result in a lot of misses:

```
int A[huge_array_size];  
for (i=0; i<some_big_N; i++) {  
    A[random(i) % huge_array_size] = i;  
}
```

where “random(i)” produces an integer random number between 0 and the maximum 32 bit integer.

QUESTION 6 (10 marks)

A machine has a physical address width of 32 bits and a virtual address width of 32 bits. The following content of a full associative 4-entry TLB:

Valid Bit	Virtual Page Number	Physical Page Number	Dirty	Ref	Access
1	0x41081	0x1111	1	20	rw
1	0x6245	0xDEED	0	1	rx
1	0x57C	0x4AC	0	8	rw
0	0x30A2	0x221	1	9	rw

- (6a) Suppose an access to the memory location at virtual address 0x000AF9F4 is translated to the physical address 0x000959F4, what is the size of a page in this machine? (2 marks)

ANSWER:

8Kbytes.

- (6b) Suppose a malicious program tries to overwrite the text segment of the code by *writing* to the virtual page 0x6245 in the hope that the program will execute it, how would the system respond? (2 marks)

ANSWER:

A segmentation fault will be raised as the page is not writable.

(6c) A file consists of 8 bytes only. The following shows its content in hexadecimal.

File byte position	Byte content
0	0xD0
1	0xE1
2	0xA2
3	0xD3
4	0xB4
5	0xE5
6	0xE6
7	0xF0

Show how this file would be stored on 5 disk RAID level 3 system employing *odd* parity. “bit[*i*][*j*]” refers to “bit *j* of byte *i*” (with ‘0’ being the least significant bit/byte). You need to write down the position as well as the value in the answer template. (6 marks)

ANSWER:

Disk 0	Disk 1	Disk 2	Disk 3	Disk 4
bit[0][0] = 0	bit[0][1] = 0	bit[0][2] = 0	bit[0][3] = 0	bit = 1
bit[0][4] = 1	bit[0][5] = 0	bit[0][6] = 1	bit[0][7] = 1	bit = 0
bit[1][0] = 1	bit[1][1] = 0	bit[1][2] = 0	bit[1][3] = 0	bit = 0
bit[1][4] = 1	bit[1][5] = 1	bit[1][6] = 1	bit[1][7] = 1	bit = 1
bit[2][0] = 0	bit[2][1] = 1	bit[2][2] = 0	bit[2][3] = 0	bit = 0
bit[2][4] = 0	bit[2][5] = 1	bit[2][6] = 0	bit[2][7] = 1	bit = 1
bit[3][0] = 1	bit[3][1] = 1	bit[3][2] = 0	bit[3][3] = 0	bit = 1
bit[3][4] = 1	bit[3][5] = 0	bit[3][6] = 1	bit[3][7] = 1	bit = 0
bit[4][0] = 0	bit[4][1] = 0	bit[4][2] = 1	bit[4][3] = 0	bit = 0
bit[4][4] = 1	bit[4][5] = 1	bit[4][6] = 0	bit[4][7] = 1	bit = 0
bit[5][0] = 1	bit[5][1] = 0	bit[5][2] = 1	bit[5][3] = 0	bit = 1
bit[5][4] = 0	bit[5][5] = 1	bit[5][6] = 1	bit[5][7] = 1	bit = 0
bit[6][0] = 0	bit[6][1] = 1	bit[6][2] = 1	bit[6][3] = 0	bit = 1
bit[6][4] = 0	bit[6][5] = 1	bit[6][6] = 1	bit[6][7] = 1	bit = 0
bit[7][0] = 1	bit[7][1] = 1	bit[7][2] = 1	bit[7][3] = 0	bit = 0
bit[7][4] = 1	bit[7][5] = 1	bit[7][6] = 1	bit[7][7] = 1	bit = 1

=== END OF PAPER ===